

Problem 1

We assume that the feed-forward neural network has L hidden layers, where each layer applies a linear activation function.

Assume the input is x , then we have:

$$h^{(1)} = W_1 x + b_1$$

$$h^{(2)} = W_2 h^{(1)} + b_2 = W_2 W_1 x + W_2 b_1 + b_2$$

$$\begin{aligned} h^{(3)} &= W_3 h^{(2)} + b_3 = W_3 (W_2 W_1 x + W_2 b_1 + b_2) + b_3 \\ &= W_3 W_2 W_1 x + \underbrace{W_3 W_2 b_1 + W_3 b_2 + b_3}_{\text{bias terms}} \end{aligned}$$

$$\vdots$$

$$\text{Then, } h^{(L)} = W_L W_{L-1} \cdots W_1 x + \text{bias terms}$$

$$= W_{\text{combined}} \cdot x + b_{\text{combined}}$$

where $W_{\text{combined}} = W_L W_{L-1} \cdots W_1$ and b_{combined} is a combination of weighted bias.

Therefore, even though we have multiple layers, since the activation function is linear, the network behaves exactly like a single linear transformation with no hidden layer.


Problem 2

Description: For this question, I extended the 2-layer logistic regression to 6 architectures stated in the question. The neuron choices are 50 or 100. The optimizers are Adam or RMSProp. Then I plotted the loss & accuracy curves as well as linear interpolation plot for each setting. Finally, I discussed the impact of different hyperparameters on learning rates and activation functions.

Results

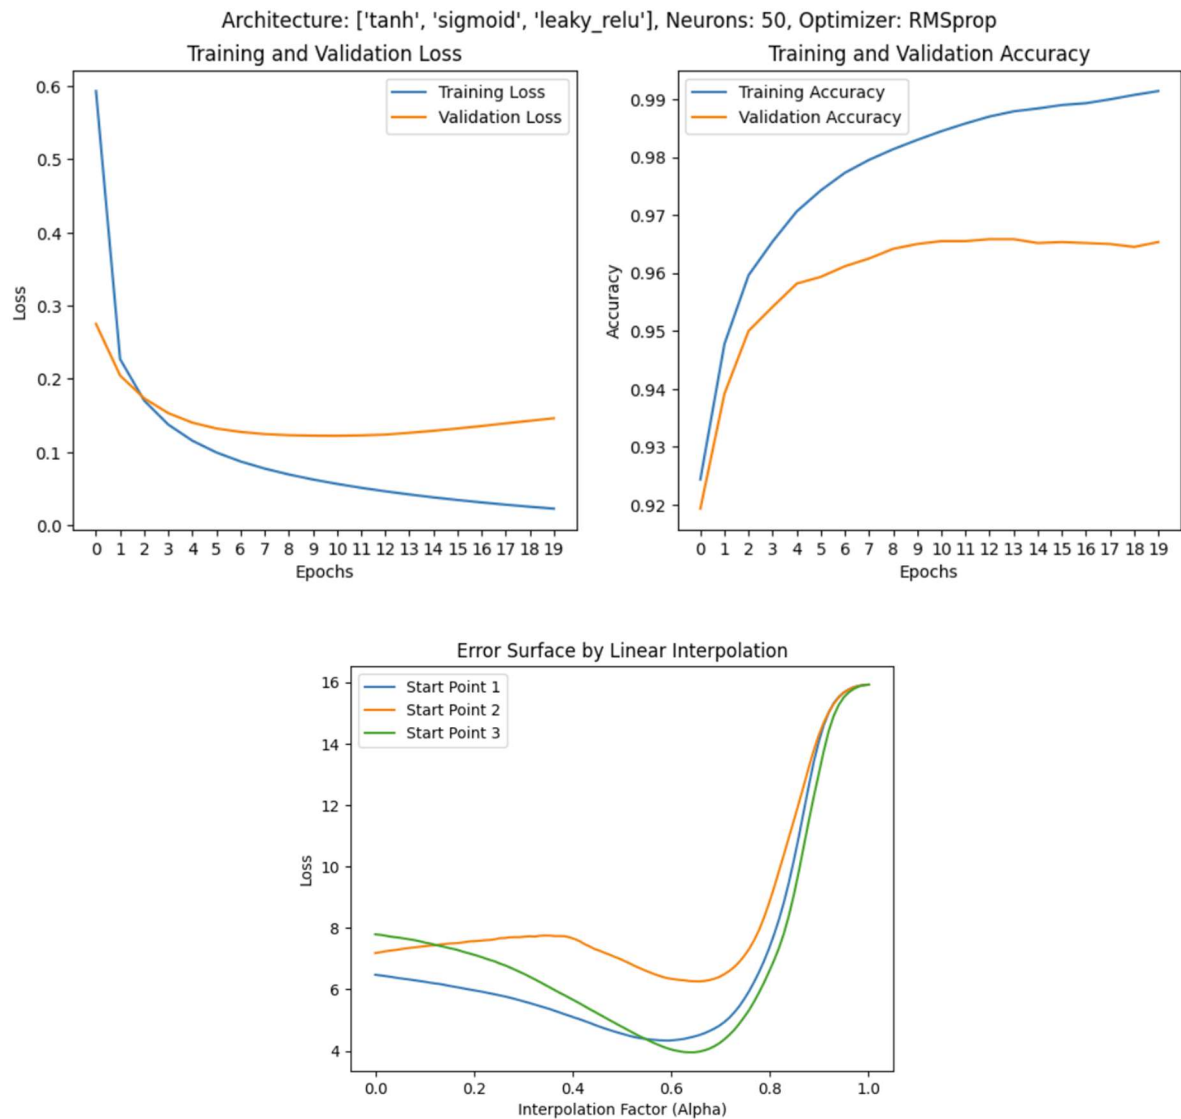
Architecture 1: 1st layer: tanh, 2nd layer: sigmoid, 3rd layer: leaky ReLU

Adam as optimizer with 50 neurons:




- Final Validation Accuracy: 96.54%

RMSProp as optimizer with 50 neurons:



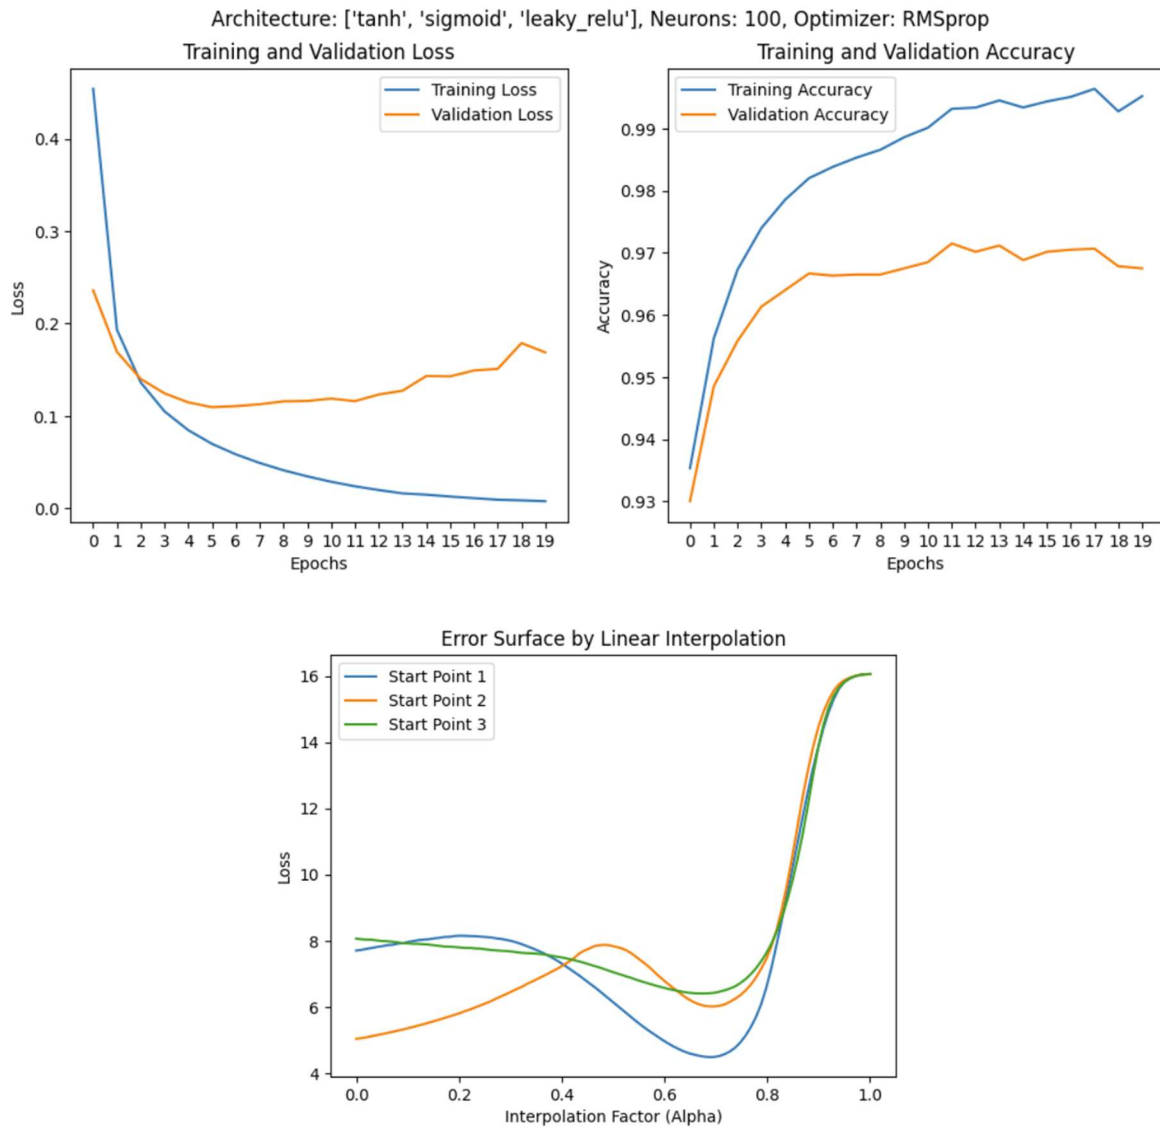
- Final Validation Accuracy: 96.86%

Adam as optimizer with 100 neurons:



- Final Validation Accuracy: 97.10%


RMSProp as optimizer with 100 neurons:



- Final Validation Accuracy: 97.28%

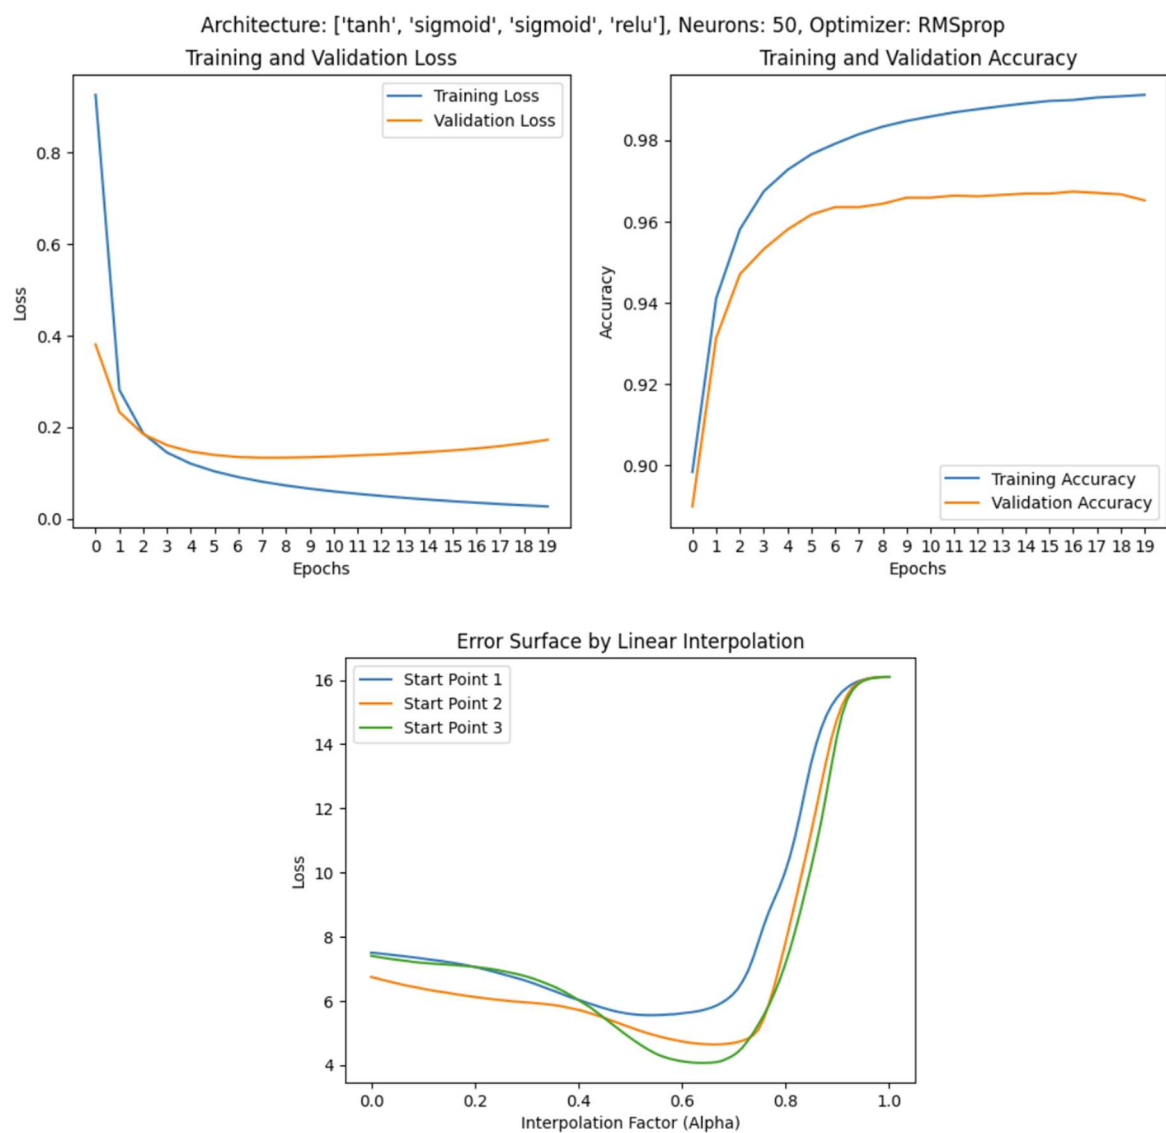
Architecture 2: 1st layer: tanh, 2nd layer: sigmoid, 3rd layer: sigmoid, 4th layer: ReLU

Adam as optimizer with 50 neurons:



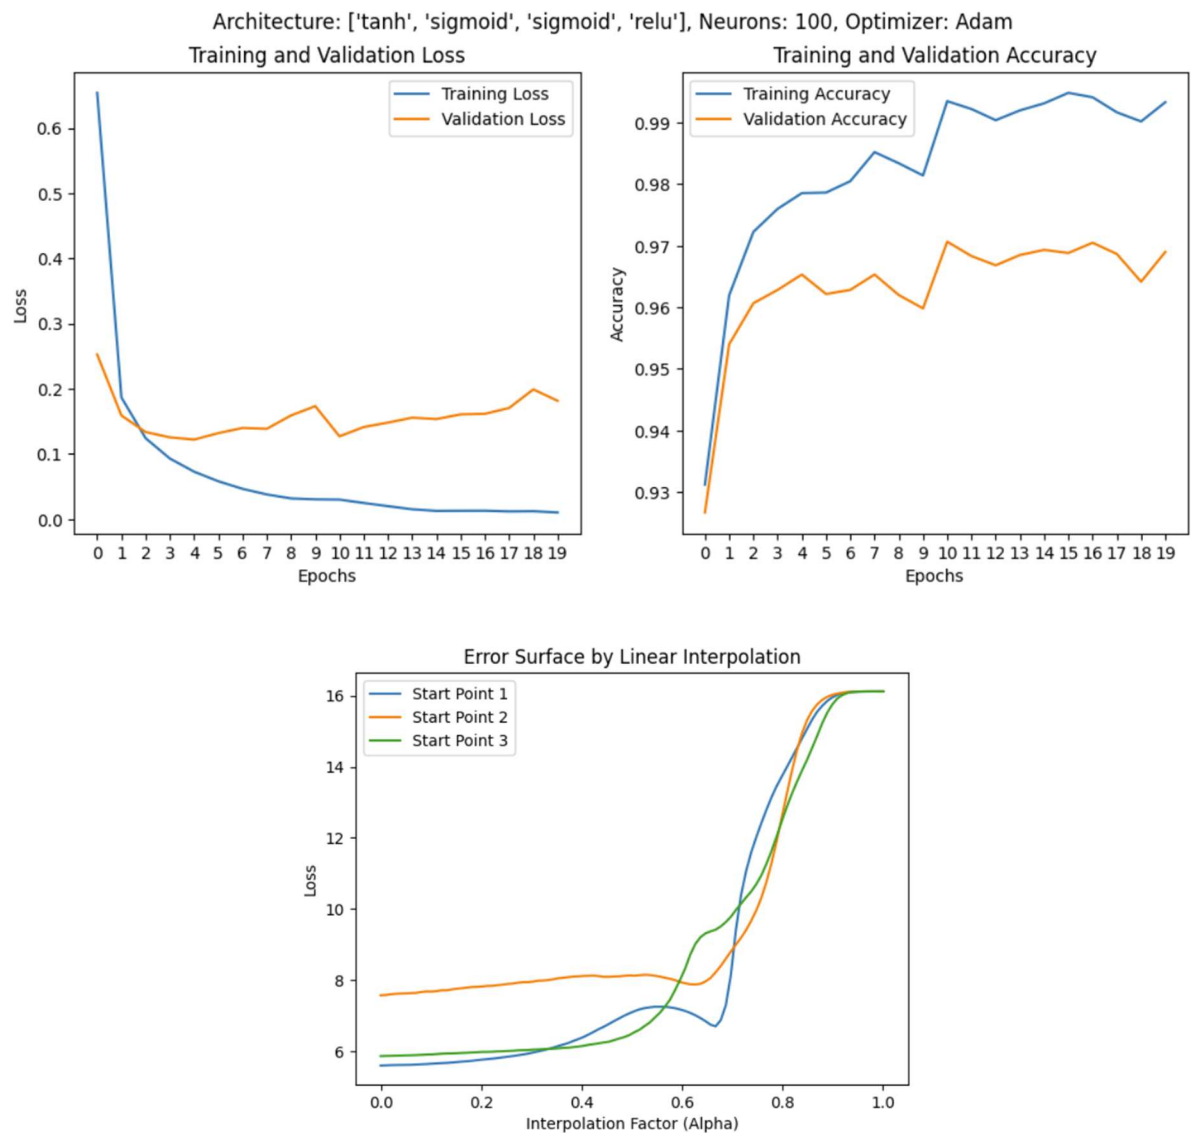
- Final Validation Accuracy: 96.52%

RMSProp as optimizer with 50 neurons:



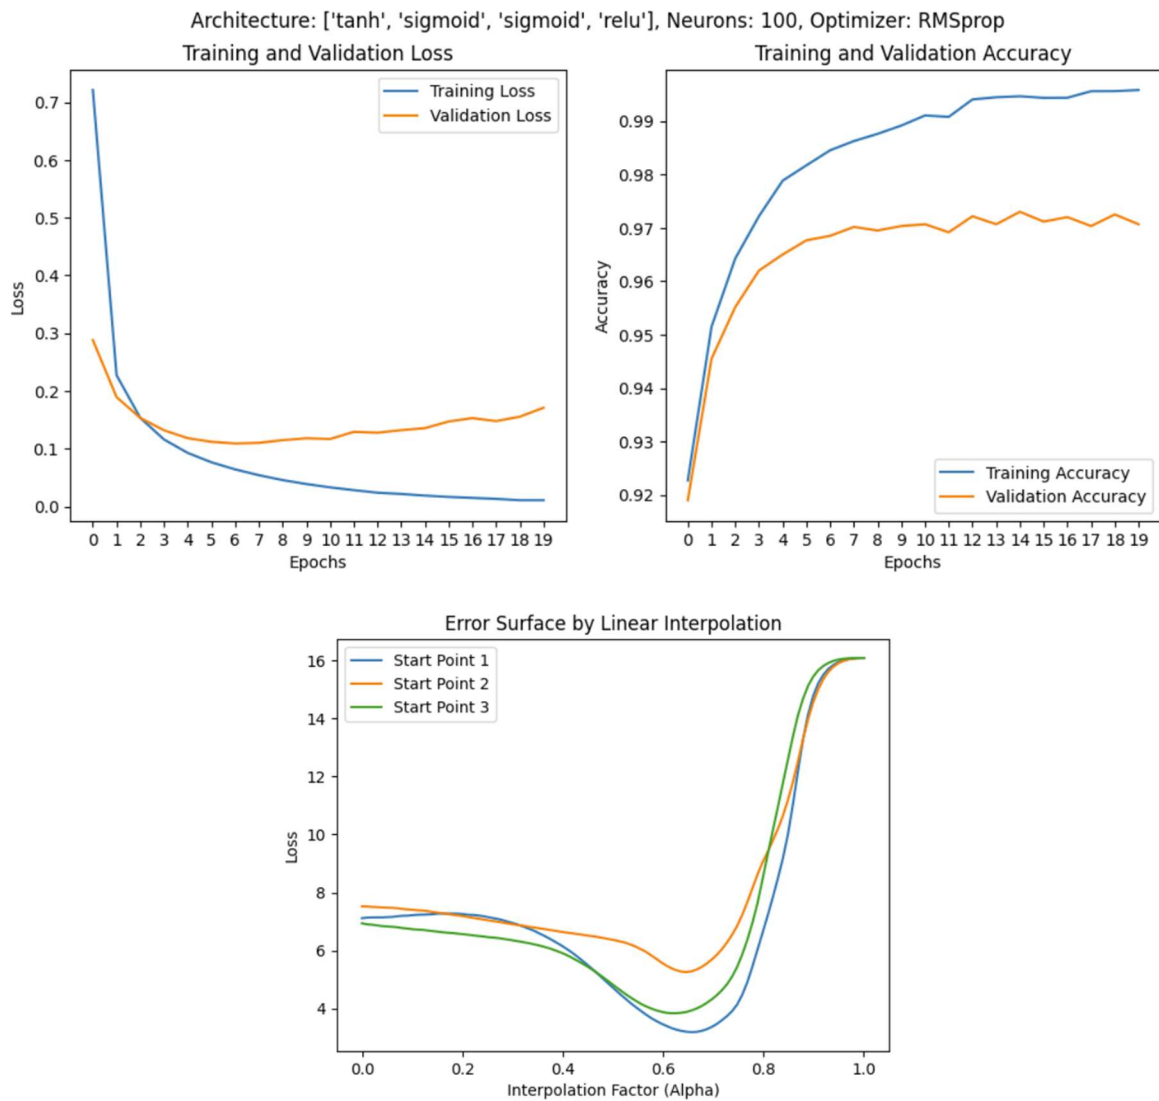
- Final Validation Accuracy: 96.71%

Adam as optimizer with 100 neurons:



- Final Validation Accuracy: 96.95%

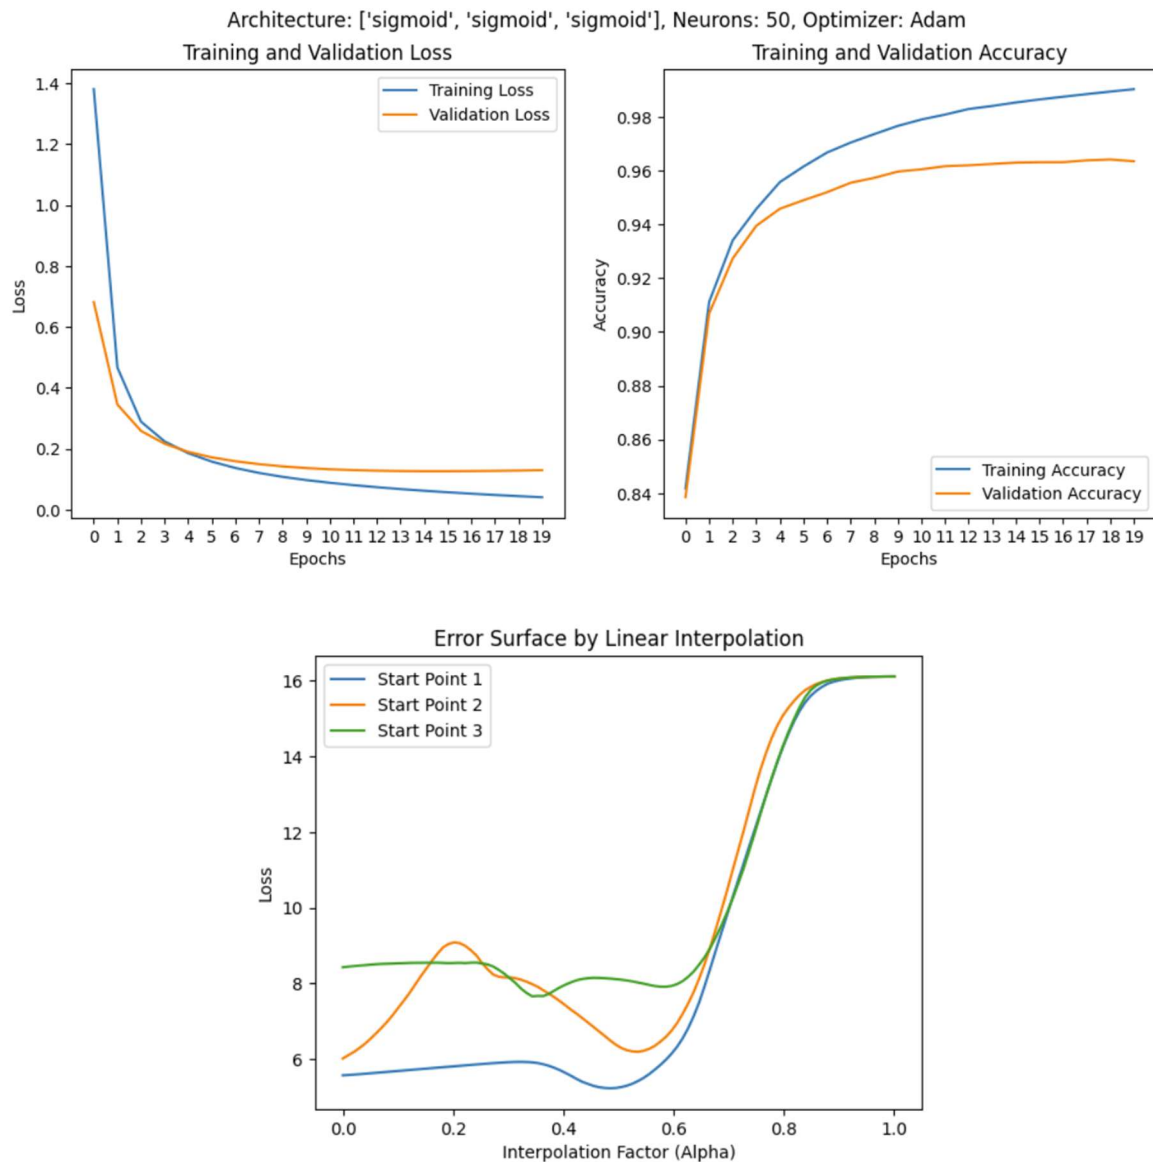
RMSProp as optimizer with 100 neurons:



- Final Validation Accuracy: 97.36%


Architecture 3: 3 layers, all sigmoid

Adam as optimizer with 50 neurons:



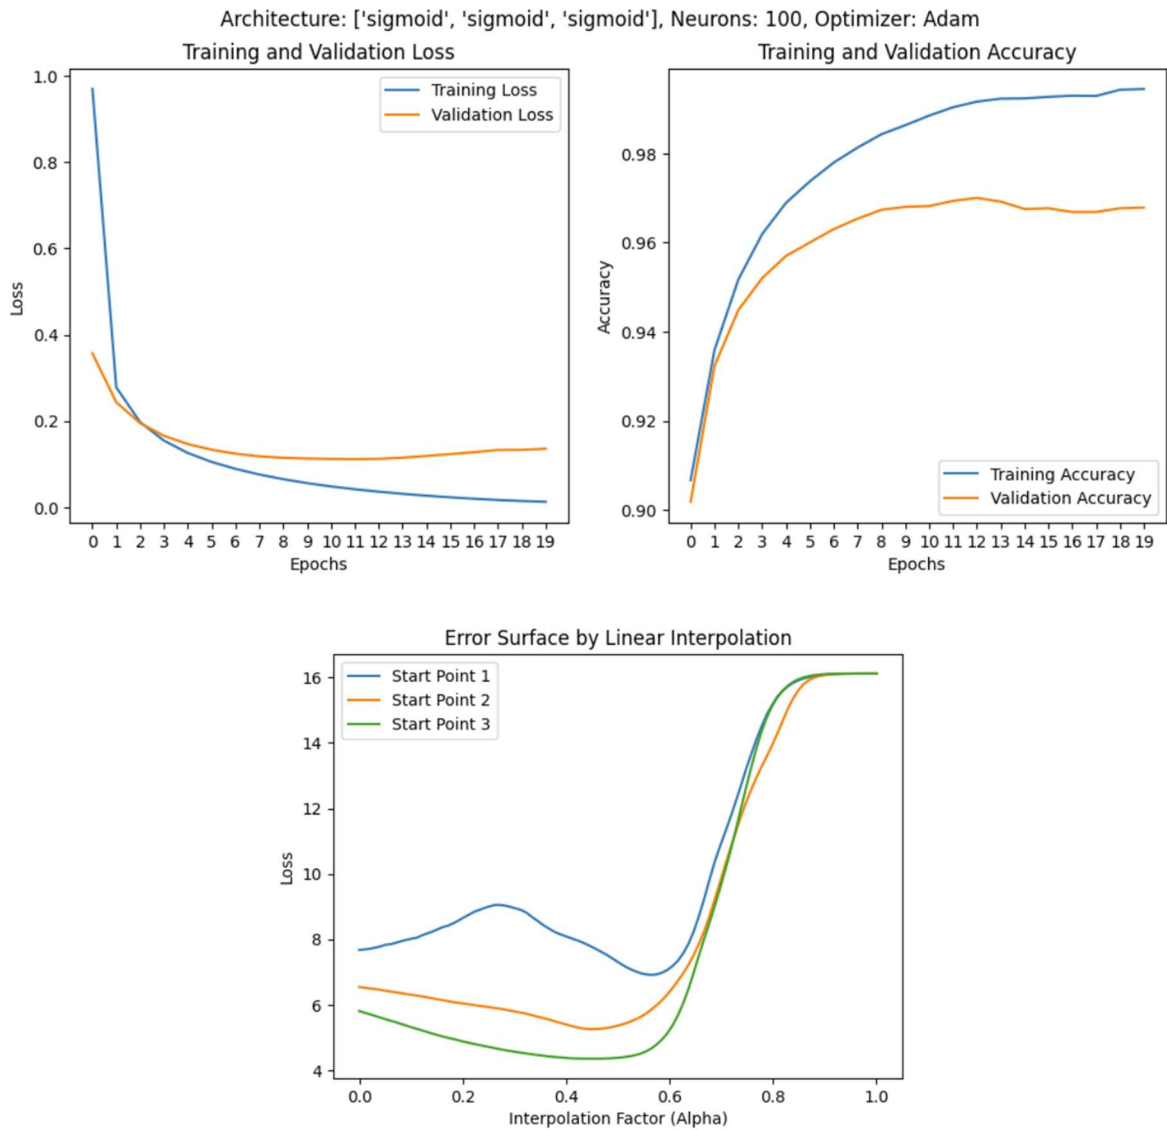
- Final Validation Accuracy: 96.90%

RMSProp as optimizer with 50 neurons:



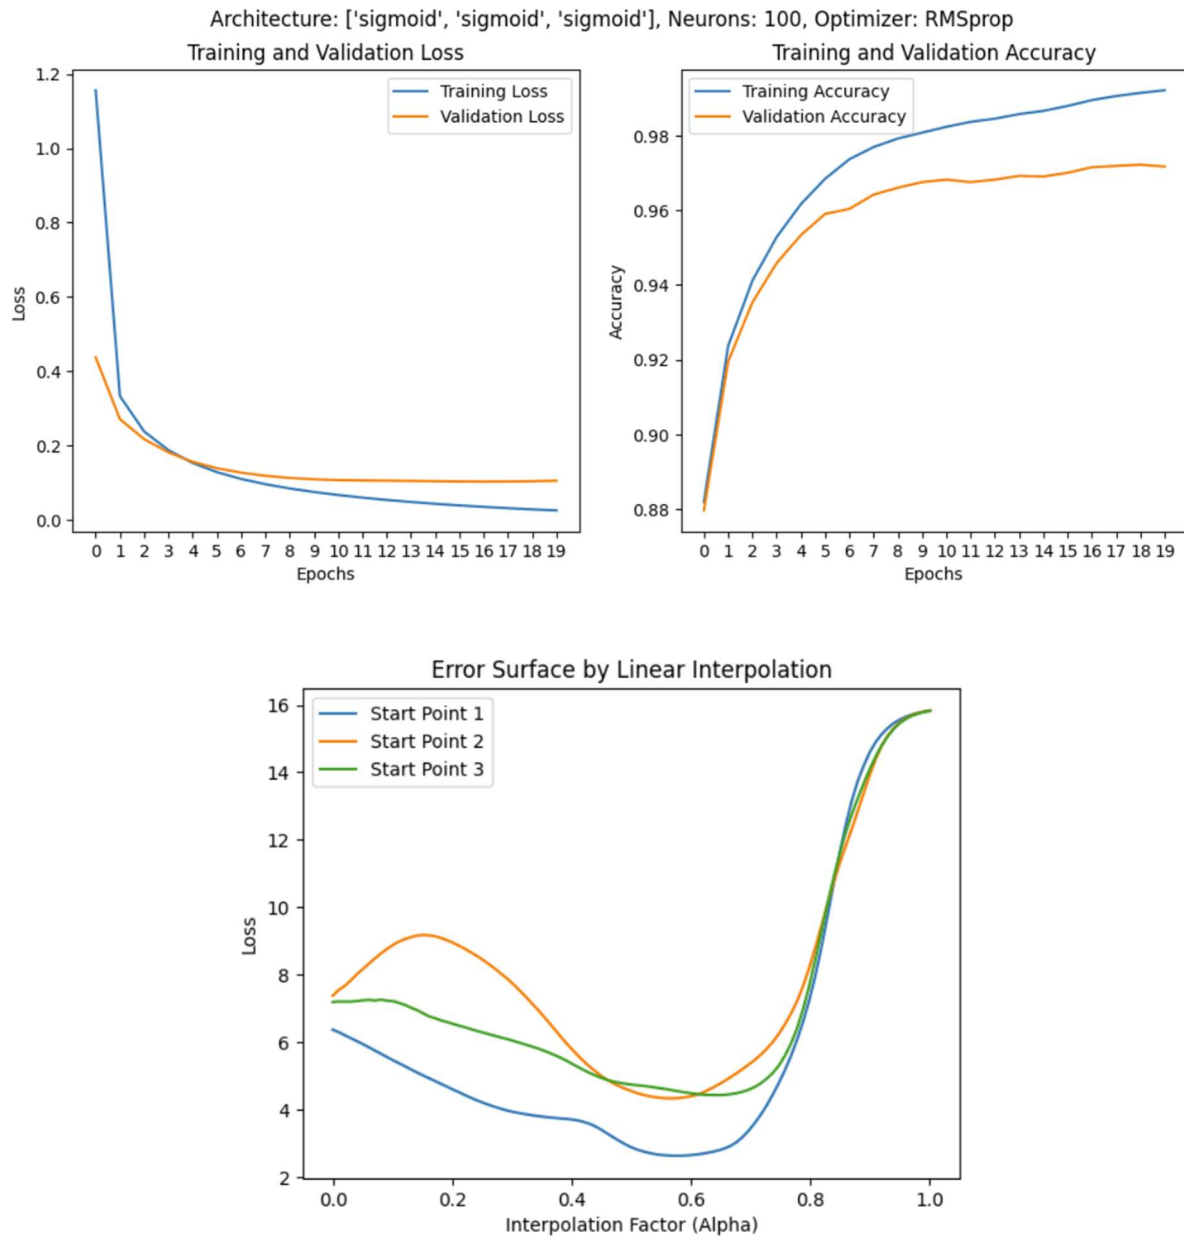
- Final Validation Accuracy: 96.84%

Adam as optimizer with 100 neurons:



- Final Validation Accuracy: 97.25%


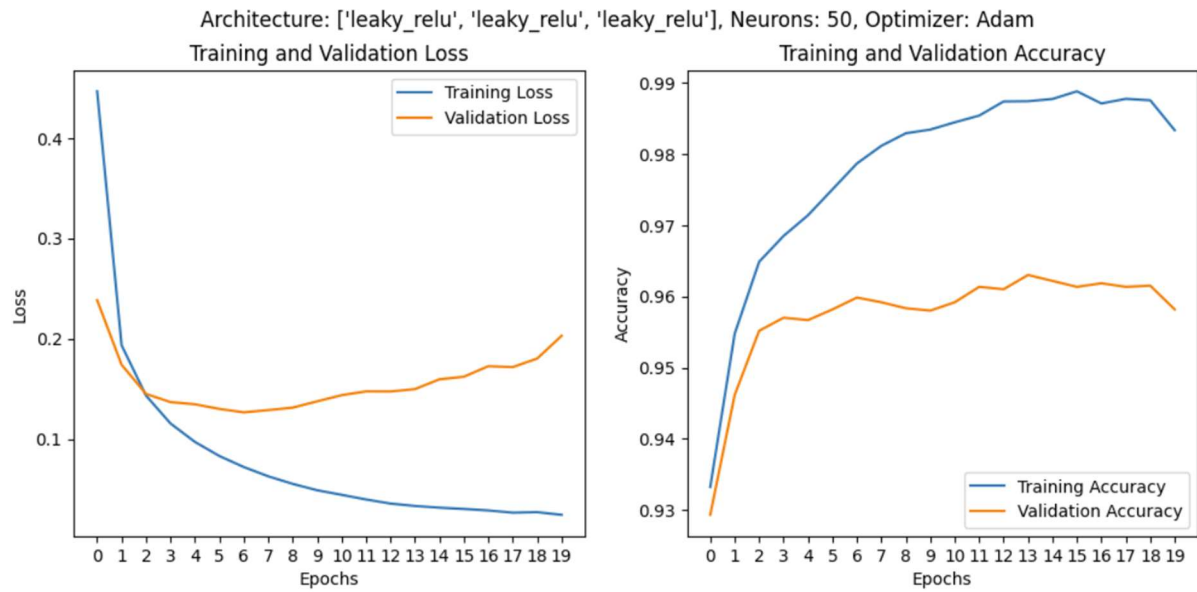
RMSProp as optimizer with 100 neurons:



- Final Validation Accuracy: 97.63%

Architecture 4: 3 layers, all leaky ReLU

Adam as optimizer with 50 neurons:




- Final Validation Accuracy: 96.45%


RMSProp as optimizer with 50 neurons:

Architecture: ['leaky_relu', 'leaky_relu', 'leaky_relu'], Neurons: 50, Optimizer: RMSprop


Training and Validation Loss



Training and Validation Accuracy



Error Surface by Linear Interpolation




- Final Validation Accuracy: 96.97%

Adam as optimizer with 100 neurons:

Architecture: ['leaky_relu', 'leaky_relu', 'leaky_relu'], Neurons: 100, Optimizer: Adam

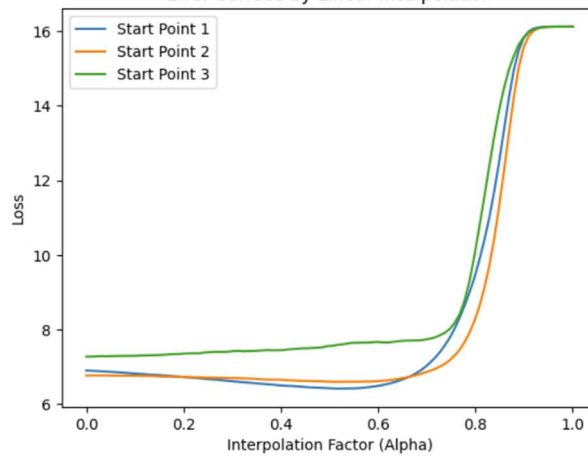
Training and Validation Loss



Training and Validation Accuracy



Error Surface by Linear Interpolation




- Final Validation Accuracy: 97.55%


RMSProp as optimizer with 100 neurons:


Architecture: ['leaky_relu', 'leaky_relu', 'leaky_relu'], Neurons: 100, Optimizer: RMSprop

Training and Validation Loss



Training and Validation Accuracy






- Final Validation Accuracy: 97.75%

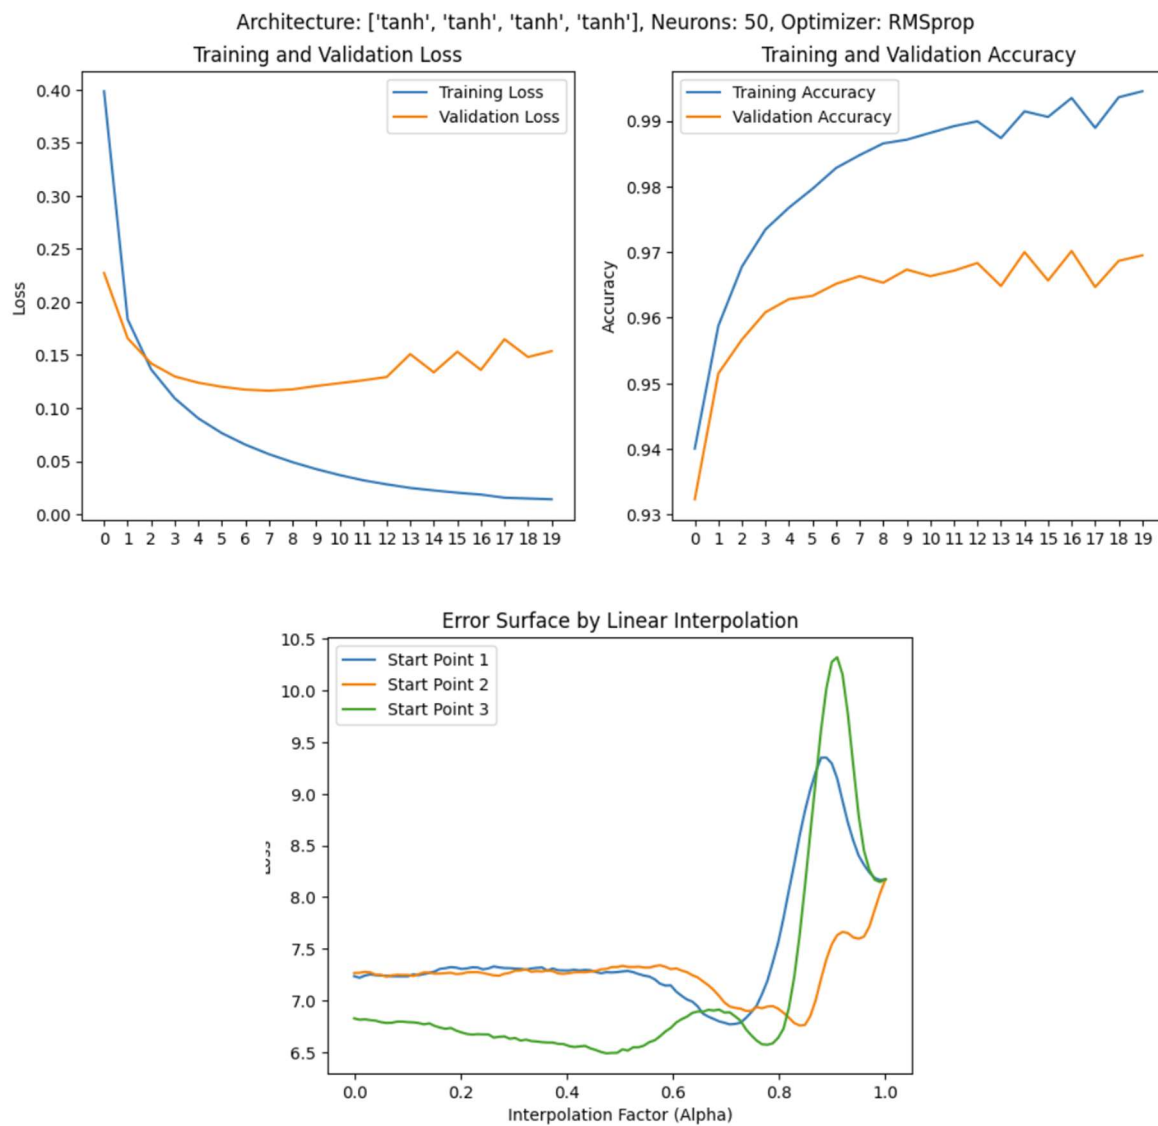
Architecture 5: 4 layers, all tanh

Adam as optimizer with 50 neurons:




- Final Validation Accuracy: 96.77%

RMSProp as optimizer with 50 neurons:



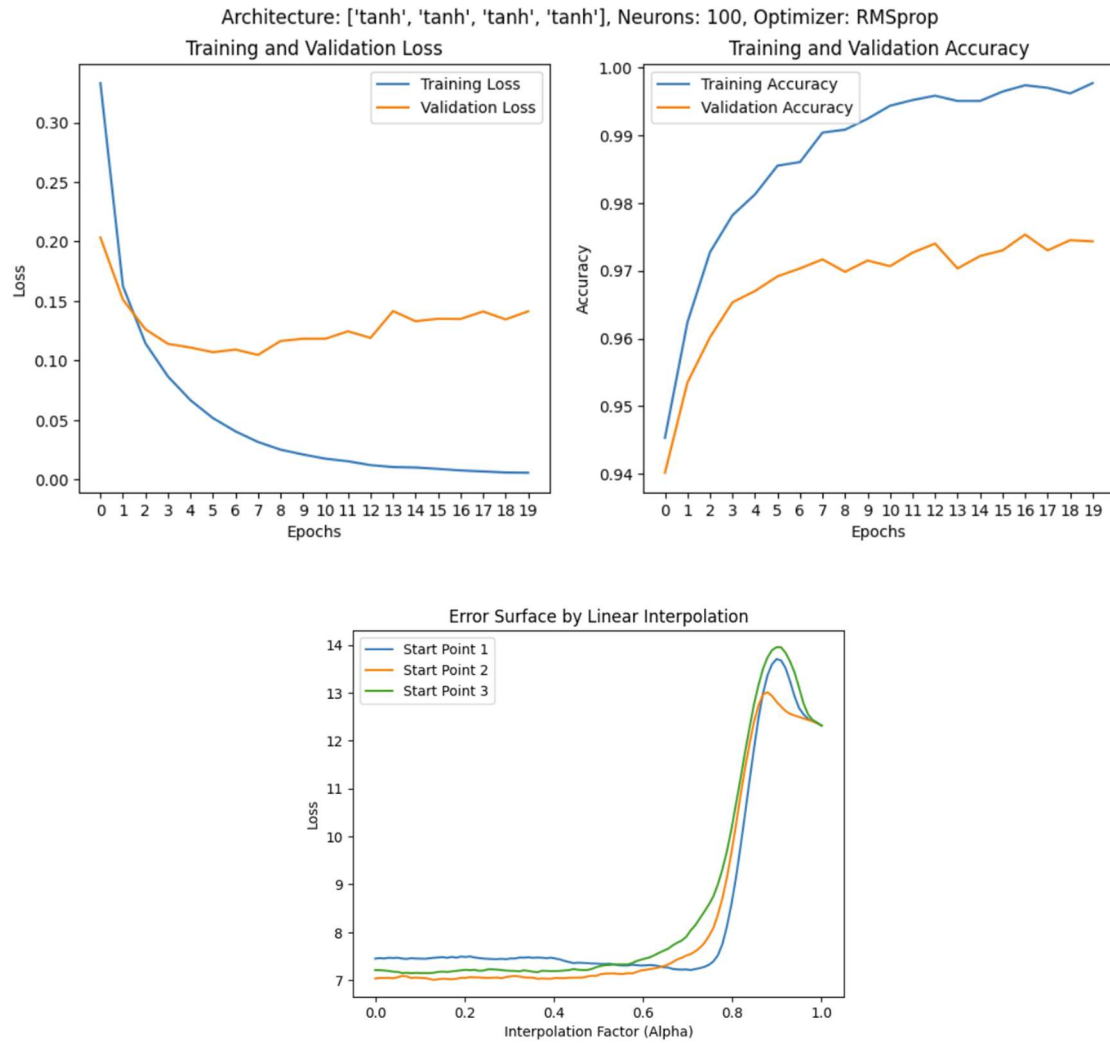
- Final Validation Accuracy: 97.19%

Adam as optimizer with 100 neurons:



- Final Validation Accuracy: 97.57%


RMSProp as optimizer with 100 neurons:



- Final Validation Accuracy: 97.32%



Architecture 6: 4 layers, all leaky ReLU

Adam as optimizer with 50 neurons:




- Final Validation Accuracy: 96.27%


RMSProp as optimizer with 50 neurons:



- Final Validation Accuracy: 96.95%



Adam as optimizer with 100 neurons:





- Final Validation Accuracy: 97.57%

RMSProp as optimizer with 100 neurons:



- Final Validation Accuracy: 97.78%

About the Error Surface

For the 2D error surface by linear interpolation, we interpolated between a few different randomly initialized starting weights and the final optimized weight to visualize the error surface.

Note: For the linear interpolation plots, I use the formula below:

$$interpolated\ weights = (1 - \alpha) * start\ weight + \alpha * final\ weight$$

Overall Accuracy Results on Validation data

Architecture	50 Neurons (Adam/RMSProp)	100 Neurons (Adam/RMSProp)
1	0.9654 / 0.9686	0.9710 / 0.9728
2	0.9652 / 0.9671	0.9695 / 0.9736
3	0.9690 / 0.9684	0.9725 / 0.9763
4	0.9645 / 0.9697	0.9755 / 0.9775
5	0.9677 / 0.9719	0.9757 / 0.9732
6	0.9627 / 0.9695	0.9757 / 0.9778

Conclusion & Analysis

1. Impact of Neurons per Layer: Increasing the number of neurons from 50 to 100 consistently improved the validation accuracy across all architectures. This suggests that higher neuron counts allow the model to capture more complex patterns in the data, leading to better generalization.
2. Optimizer Performance: The RMSProp optimizer slightly outperformed Adam in most configurations, particularly with higher neuron counts (100). This indicates that RMSProp may be more effective for these architectures and the problem dataset, potentially due to better handling of adaptive learning rates in layered structures.
3. Architecture and Activation Function Influence: Among the architectures tested, the 4-layer tanh (architecture 5) and 4-layer leaky ReLU (architecture 6) achieved the highest validation accuracy with 100 neurons, indicating that deeper architectures with non-linear activations (tanh, leaky ReLU) help in capturing complex non-linear relationships, which may contribute to better overall model performance.

Different Learning Rate

For this part, I choose the architecture 1 with 50 neurons and Adam as optimizer to train with different learning rates [0.0001, 0.001, 0.01, 0.1, 0.5]. We want to discover how the learning rate affects the training performance over time.

Result

Learning rate	Validation accuracy	Testing accuracy	Running time (in second)
0.0001	0.9523	0.9506	460.09
0.001	0.9648	0.9648	456.17
0.01	0.9624	0.9642	436.34
0.1	0.1141	0.1134	434.87
0.5	0.1141	0.1134	432.27

Conclusion

1. Optimal Learning Rate Range: Learning rates of 0.001 and 0.01 achieved the highest validation and testing accuracy, suggesting that these values strike a good balance between convergence speed and stability for this architecture and optimizer combination.
2. Effect of Very Low Learning Rate (0.0001): While the lowest learning rate (0.0001) provided relatively high accuracy, it took the longest time to converge, indicating that very low learning rates can increase training duration without significant accuracy benefits over slightly higher learning rates.
3. Instability with High Learning Rates (0.1 and 0.5): Both the 0.1 and 0.5 learning rates resulted in significantly lower accuracy scores, which indicates that these high values lead to unstable training, potentially causing the model to diverge or fail to find an optimal solution.
4. Marginal Time Differences Across Learning Rates: The running time slightly decreased as the learning rate increased, with only marginal differences across learning rates, suggesting that higher learning rates reduce the number of required iterations but may compromise accuracy if set too high.

code.py

```
1 import tensorflow as tf
2 print(tf.__version__)
3
4 import os
5 import time
6
7 import numpy as np # linear algebra
8 import matplotlib.pyplot as plt
9 from tensorflow.keras.layers import Dense, LeakyReLU
10 from tensorflow.keras.optimizers import Adam, RMSprop
11 import warnings
12 warnings.filterwarnings('ignore')
13
14
15 # generate original training and test data
16 img_size = 28
17 n_classes = 10
18
19 #MNIST data image of shape 28*28=784
20 input_size = 784
21
22 # 0-9 digits recognition (labels)
23 output_size = 10
24
25 #-----
26 #option 1: load MNIST dataset
27 #from tensorflow.examples.tutorials.mnist import input_data
28 #mnist = input_data.read_data_sets("data/", one_hot=True)
29
30
31 #-----
32 #option 2: load MNIST dataset
33 print('\nLoading MNIST')
34 mnist = tf.keras.datasets.mnist
35 (x_train, y_train), (x_test, y_test) = mnist.load_data()
36
37 x_train = np.reshape(x_train, [-1, img_size*img_size])
38 x_train = x_train.astype(np.float32)/255
39
40 x_test = np.reshape(x_test, [-1, img_size*img_size])
41 x_test = x_test.astype(np.float32)/255
42
43 to_categorical = tf.keras.utils.to_categorical
44 y_train = to_categorical(y_train)
45 y_test = to_categorical(y_test)
46
47 print('\nSplitting data')
48
49 ind = np.random.permutation(x_train.shape[0])
50 x_train, y_train = x_train[ind], y_train[ind]
51
```

```

52 # 10% for validation
53 validationPct = 0.1
54 n = int(x_train.shape[0] * (1-validationPct))
55 x_valid = x_train[n:]
56 x_train = x_train[:n]
57 #
58 y_valid = y_train[n:]
59 y_train = y_train[:n]
60
61 train_num_examples = x_train.shape[0]
62 valid_num_examples = x_valid.shape[0]
63 test_num_examples = x_test.shape[0]
64
65 print(train_num_examples, valid_num_examples, test_num_examples)
66
67 # Global Parameters
68 #-----
69 # learning rate
70 learning_rate = 0.05
71
72 #training_epochs = 1000
73 #batch_size = 30
74
75 training_epochs = 100
76 batch_size = 50
77
78 display_step = 10
79
80 #Network Architecture
81 # -----
82 #
83 # Two hidden layers
84 #
85 #-----
86 # number of neurons in layer 1
87 n_hidden_1 = 200
88 # number of neurons in layer 2
89 n_hidden_2 = 300
90
91 #MNIST data image of shape 28*28=784
92 input_size = 784
93
94 # 0-9 digits recognition (labels)
95 output_size = 10
96
97 def loss_2(output, y):
98     """
99     Computes softmax cross entropy between logits and labels and returns the loss.
100
101     Input:
102         - output: the output (logits) of the inference function (shape: batch_size *
103         num_of_classes)
104         - y: true labels for the sample batch (shape: batch_size * num_of_classes)
105     Output:

```

```

105         - loss: the scalar loss value for the batch
106     """
107     # Computes softmax cross entropy between logits (output) and true labels (y)
108     xentropy = tf.nn.softmax_cross_entropy_with_logits(logits=output, labels=y)
109
110     # Return the mean cross-entropy loss across the batch
111     loss = tf.reduce_mean(xentropy)
112
113     return loss
114
115 def evaluate(output, y):
116     """
117     Evaluates the accuracy on the validation set.
118     Input:
119         - output: prediction vector of the network for the validation set
120         - y: true value for the validation set
121     Output:
122         - accuracy: accuracy on the validation set (scalar between 0 and 1)
123     """
124     # Check if the predicted class equals the true class
125     correct_prediction = tf.equal(tf.argmax(output, 1), tf.argmax(y, 1))
126
127     # Compute accuracy as the mean of correct predictions
128     accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
129
130     # Log validation accuracy using TensorFlow summary (if needed)
131     with tf.summary.create_file_writer('./logs/validation').as_default():
132         tf.summary.scalar("validation_error", 1.0 - accuracy, step=0)
133
134     return accuracy
135
136 def build_model(architecture, neurons_per_layer):
137     model = tf.keras.Sequential()
138     input_shape = (input_size,)
139
140     for i, activation in enumerate(architecture):
141         if activation == 'leaky_relu':
142             model.add(Dense(neurons_per_layer, input_shape=input_shape if i == 0 else None))
143             model.add(LeakyReLU(negative_slope=0.01))
144         else:
145             model.add(Dense(neurons_per_layer, activation=activation,
146                             input_shape=input_shape if i == 0 else None))
147
148     model.add(Dense(output_size))
149     return model
150
151 # Function to plot training history
152 def plot_training_history(history, architecture_name, neurons_per_layer, optimizer_name,
153 epochs):
154     plt.figure(figsize=(12, 5))
155
156     # Add supertitle for the architecture configuration
157     plt.suptitle(f"Architecture: {architecture_name}, Neurons: {neurons_per_layer},
158 Optimizer: {optimizer_name}")

```

```

156 # Plot training and validation loss
157 plt.subplot(1, 2, 1)
158 plt.plot(history['loss'], label='Training Loss')
159 plt.plot(history['val_loss'], label='Validation Loss')
160 plt.xlabel('Epochs')
161 plt.ylabel('Loss')
162 plt.legend()
163 plt.title('Training and Validation Loss')
164 plt.xticks(np.arange(0, epochs, 1))
165
166 # Plot training and validation accuracy
167 plt.subplot(1, 2, 2)
168 plt.plot(history['accuracy'], label='Training Accuracy')
169 plt.plot(history['val_accuracy'], label='Validation Accuracy')
170 plt.xlabel('Epochs')
171 plt.ylabel('Accuracy')
172 plt.legend()
173 plt.title('Training and Validation Accuracy')
174 plt.xticks(np.arange(0, epochs, 1))
175
176 plt.show()
177
178 def visualize_error_surface(model, x_train, y_train, start_points):
179     """
180     Visualizes the error surface by interpolating between the starting weights and final
181     trained weights.
182
183     Parameters:
184     - model: Trained model to use for error surface visualization.
185     - x_train: Training data features.
186     - y_train: Training data labels.
187     - start_points: List of initial weights (random starting points) to interpolate from.
188     """
189     loss_function = tf.keras.losses.CategoricalCrossentropy()
190
191     # Get the final trained weights
192     final_weights = model.get_weights()
193
194     # Define the interpolation factor (alpha) values
195     alphas = np.linspace(0, 1, 100)
196
197     # Plot the error surface for each starting point
198     for idx, start_weights in enumerate(start_points):
199         # Store the losses along the interpolation path
200         losses = []
201
202         # Interpolate between the starting weights and final weights
203         for alpha in alphas:
204             # Compute interpolated weights
205             interpolated_weights = [(1 - alpha) * start + alpha * final
206                                     for start, final in zip(start_weights, final_weights)]
207
208             # Set interpolated weights in the model
209             model.set_weights(interpolated_weights)

```

```

209         # Compute loss for the interpolated model
210         y_pred = model(x_train)
211         loss = loss_function(y_train, y_pred).numpy()
212         losses.append(loss)
213
214     # Plot the losses for this interpolation path
215     plt.plot(alphas, losses, label=f"Start Point {idx + 1}")
216
217     plt.xlabel("Interpolation Factor (Alpha)")
218     plt.ylabel("Loss")
219     plt.title("Error Surface by Linear Interpolation")
220     plt.legend()
221     plt.show()
222
223 def training_testing(architecture_name, neurons_per_layer, optimizer_name):
224     start_time = time.time()
225
226     # Define inputs directly (no need for placeholders)
227     input_size = 784
228     output_size = 10
229     batch_size = 128
230     training_epochs = 20
231     display_step = 5
232
233     # Instantiate the model with the architecture parameters
234     model = build_model(architecture_name, neurons_per_layer)
235
236     # Define optimizer
237     if optimizer_name == 'Adam':
238         optimizer = tf.optimizers.Adam()
239     elif optimizer_name == 'RMSprop':
240         optimizer = tf.optimizers.RMSprop()
241
242     # Define the checkpoint manager
243     checkpoint = tf.train.Checkpoint(optimizer=optimizer, model=model)
244     checkpoint_manager = tf.train.CheckpointManager(checkpoint, './logs/multi_layer',
max_to_keep=5)
245
246     # Training loop
247     history = {'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': []}
248
249     for epoch in range(training_epochs):
250         avg_cost = 0.
251         total_batch = int((train_num_examples + batch_size - 1) / batch_size)
252
253         for i in range(total_batch):
254             start = i * batch_size
255             end = min(train_num_examples, start + batch_size)
256             minibatch_x = x_train[start:end]
257             minibatch_y = y_train[start:end]
258
259             # Define training step using GradientTape
260             with tf.GradientTape() as tape:
261                 output = model(minibatch_x)

```

```

262         cost = loss_2(output, minibatch_y)
263
264         # Compute gradients and apply them
265         gradients = tape.gradient(cost, model.trainable_variables)
266         optimizer.apply_gradients(zip(gradients, model.trainable_variables))
267
268         avg_cost += cost.numpy() / total_batch
269
270         # Append metrics for plotting
271         history['loss'].append(avg_cost)
272         accuracy = evaluate(model(x_train), y_train)
273         history['accuracy'].append(accuracy)
274
275         val_loss = loss_2(model(x_valid), y_valid).numpy()
276         val_accuracy = evaluate(model(x_valid), y_valid)
277         history['val_loss'].append(val_loss)
278         history['val_accuracy'].append(val_accuracy)
279
280         if (epoch+1) % display_step == 0:
281             print(f"Epoch: {(epoch+1):2d}, cost={avg_cost:.7f}, Validation Error={1-
accuracy:.7f}, Training Accuracy={accuracy}, Validation Accuracy={val_accuracy}")
282             checkpoint_manager.save()
283
284         # Final test accuracy
285         accuracy = evaluate(model(x_test), y_test)
286         print("Test Accuracy:", accuracy)
287
288         elapsed_time = time.time() - start_time
289         print(f'Execution time (seconds) was {elapsed_time:.3f}')
290
291         # Call the plotting function to visualize training history
292         plot_training_history(history, architecture_name, neurons_per_layer, optimizer_name,
training_epochs)
293
294         # Generate a few random starting points for error surface visualization
295         start_points = [ [np.random.normal(size=w.shape) for w in model.get_weights()] for _ in
range(3)]
296
297         # Visualize error surface by linear interpolation
298         visualize_error_surface(model, x_train, y_train, start_points)
299
300     def training_testing_diff_lr(architecture_name, neurons_per_layer, optimizer_name,
learning_rates):
301         results = [] # Store validation accuracy and elapsed time for each learning rate
302
303         # Define inputs
304         input_size = 784
305         output_size = 10
306         batch_size = 128
307         training_epochs = 20
308         display_step = 5
309
310         # Loop through different learning rates
311         for lr in learning_rates:

```

```

312     print(f"\nTraining with learning rate: {lr}")
313
314     # Start timer for this learning rate
315     start_time = time.time()
316
317     # Instantiate the model with the architecture parameters
318     model = build_model(architecture_name, neurons_per_layer)
319
320     # Define optimizer with the specific learning rate
321     if optimizer_name == 'Adam':
322         optimizer = tf.optimizers.Adam(learning_rate=lr)
323     elif optimizer_name == 'RMSprop':
324         optimizer = tf.optimizers.RMSprop(learning_rate=lr)
325
326     # Define the checkpoint manager
327     checkpoint = tf.train.Checkpoint(optimizer=optimizer, model=model)
328     checkpoint_manager = tf.train.CheckpointManager(checkpoint, './logs/multi_layer',
max_to_keep=5)
329
330     # Training loop
331     history = {'loss': [], 'accuracy': [], 'val_loss': [], 'val_accuracy': []}
332
333     for epoch in range(training_epochs):
334         avg_cost = 0.
335         total_batch = int((train_num_examples + batch_size - 1) / batch_size)
336
337         for i in range(total_batch):
338             start = i * batch_size
339             end = min(train_num_examples, start + batch_size)
340             minibatch_x = x_train[start:end]
341             minibatch_y = y_train[start:end]
342
343             # Define training step using GradientTape
344             with tf.GradientTape() as tape:
345                 output = model(minibatch_x)
346                 cost = loss_2(output, minibatch_y)
347
348             # Compute gradients and apply them
349             gradients = tape.gradient(cost, model.trainable_variables)
350             optimizer.apply_gradients(zip(gradients, model.trainable_variables))
351
352             avg_cost += cost.numpy() / total_batch
353
354         # Append metrics for plotting
355         history['loss'].append(avg_cost)
356         accuracy = evaluate(model(x_train), y_train)
357         history['accuracy'].append(accuracy)
358
359         val_loss = loss_2(model(x_valid), y_valid).numpy()
360         val_accuracy = evaluate(model(x_valid), y_valid)
361         history['val_loss'].append(val_loss)
362         history['val_accuracy'].append(val_accuracy)
363
364         if (epoch+1) % display_step == 0:

```



```

365         print(f"Epoch: {(epoch+1):2d}, cost={avg_cost:.7f}, Validation Error={1-
accuracy:.7f}, Training Accuracy={accuracy}, Validation Accuracy={val_accuracy}")
366         checkpoint_manager.save()
367
368     # Final test accuracy
369     test_accuracy = evaluate(model(x_test), y_test)
370     print("Test Accuracy:", test_accuracy)
371
372     # Record elapsed time
373     elapsed_time = time.time() - start_time
374     print(f'Execution time (seconds) was {elapsed_time:.3f}')
375
376     # Store results for this learning rate
377     results.append({
378         'learning_rate': lr,
379         'final_val_accuracy': val_accuracy,
380         'test_accuracy': test_accuracy,
381         'time': elapsed_time
382     })
383
384     # Display summary of results for all learning rates
385     for result in results:
386         print(f"Learning Rate: {result['learning_rate']}, Final Validation Accuracy:
{result['final_val_accuracy']}, Test Accuracy: {result['test_accuracy']}, Time Taken:
{result['time']} seconds")
387
388     # Return results for further analysis if needed
389     return results
390
391 if __name__ == '__main__':
392     architectures = {
393         "1_tanh_2_sigmoid_3_leaky_relu": ['tanh', 'sigmoid', 'leaky_relu'],
394         "1_tanh_2_sigmoid_3_sigmoid_4_relu": ['tanh', 'sigmoid', 'sigmoid', 'relu'],
395         "3_layers_sigmoid": ['sigmoid', 'sigmoid', 'sigmoid'],
396         "3_layers_leaky_relu": ['leaky_relu', 'leaky_relu', 'leaky_relu'],
397         "4_layers_tanh": ['tanh', 'tanh', 'tanh', 'tanh'],
398         "4_layers_leaky_relu": ['leaky_relu', 'leaky_relu', 'leaky_relu', 'leaky_relu']}
399
400     experiment_configurations = [
401         ("1_tanh_2_sigmoid_3_leaky_relu", 50, "Adam"),
402         ("1_tanh_2_sigmoid_3_leaky_relu", 50, "RMSprop"),
403         ("1_tanh_2_sigmoid_3_leaky_relu", 100, "Adam"),
404         ("1_tanh_2_sigmoid_3_leaky_relu", 100, "RMSprop"),
405         ("1_tanh_2_sigmoid_3_sigmoid_4_relu", 50, "Adam"),
406         ("1_tanh_2_sigmoid_3_sigmoid_4_relu", 50, "RMSprop"),
407         ("1_tanh_2_sigmoid_3_sigmoid_4_relu", 100, "Adam"),
408         ("1_tanh_2_sigmoid_3_sigmoid_4_relu", 100, "RMSprop"),
409         ("3_layers_sigmoid", 50, "Adam"),
410         ("3_layers_sigmoid", 50, "RMSprop"),
411         ("3_layers_sigmoid", 100, "Adam"),
412         ("3_layers_sigmoid", 100, "RMSprop"),
413         ("3_layers_leaky_relu", 50, "Adam"),
414         ("3_layers_leaky_relu", 50, "RMSprop"),
415         ("3_layers_leaky_relu", 100, "Adam"),

```

```
416     ("3_layers_leaky_relu", 100, "RMSprop"),
417     ("4_layers_tanh", 50, "Adam"),
418     ("4_layers_tanh", 50, "RMSprop"),
419     ("4_layers_tanh", 100, "Adam"),
420     ("4_layers_tanh", 100, "RMSprop"),
421     ("4_layers_leaky_relu", 50, "Adam"),
422     ("4_layers_leaky_relu", 50, "RMSprop"),
423     ("4_layers_leaky_relu", 100, "Adam"),
424     ("4_layers_leaky_relu", 100, "RMSprop")]
425
426     for configuration in experiment_configurations:
427         architecture_name, neurons_per_layer, optimizer_name = configuration
428         training_testing(architectures[architecture_name], neurons_per_layer,
429                           optimizer_name)
429
430     learning_rates = [0.0001, 0.001, 0.01, 0.1, 0.5]
431     results = training_testing_diff_lr(['tanh', 'sigmoid', 'leaky_relu'], 50, "Adam",
432                                       learning_rates)
```