

# Code Appendix

Name: Mark Ma

UNI: km4054

Nov 7, 2024

```
[17]: from tensorflow.keras import layers, models
      from tensorflow.keras.datasets import cifar10
      from tensorflow.keras.utils import to_categorical
      from sklearn.metrics import confusion_matrix, classification_report
      import numpy as np
      import seaborn as sns
      import matplotlib.pyplot as plt
```

## 1 Problem 1

```
[18]: # Load and preprocess CIFAR-10 data
      (x_train, y_train), (x_test, y_test) = cifar10.load_data()

      # Normalize the pixel values to be between 0 and 1
      x_train, x_test = x_train.astype('float32') / 255.0, x_test.astype('float32') / 255.0
      y_train = to_categorical(y_train, 10)
      y_test = to_categorical(y_test, 10)

      # Define the original CNN model
      def create_cnn_model():
          model = models.Sequential()

          # First convolutional block
          model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same', input_shape=(32, 32, 3)))
          model.add(layers.MaxPooling2D((2, 2)))

          # Second convolutional block
          model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
          model.add(layers.MaxPooling2D((2, 2)))

          # Third convolutional block
          model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
          model.add(layers.MaxPooling2D((2, 2)))

          # Flatten and add dense layers
```

```

model.add(layers.Flatten())
model.add(layers.Dense(128, activation='relu'))
model.add(layers.Dense(64, activation='relu'))

# Output layer
model.add(layers.Dense(10, activation='softmax')) # 10 classes for CIFAR-10

return model

# Define the modified model with an additional convolutional layer
def create_modified_cnn_model():
    model = models.Sequential()

    # First convolutional block
    model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same',
↪input_shape=(32, 32, 3)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Second convolutional block
    model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling2D((2, 2)))

    # Third convolutional block
    model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling2D((2, 2)))

    # Additional convolutional block (added layer)
    model.add(layers.Conv2D(256, (3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling2D((2, 2)))

    # Flatten and add dense layers
    model.add(layers.Flatten())
    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.Dense(64, activation='relu'))

    # Output layer
    model.add(layers.Dense(10, activation='softmax')) # 10 classes for CIFAR-10

    return model

```

```

[19]: # Instantiate both models
original_model = create_cnn_model()
modified_model = create_modified_cnn_model()

# Compile both models
original_model.compile(optimizer='adam', loss='categorical_crossentropy',
↪metrics=['accuracy'])

```

```

modified_model.compile(optimizer='adam', loss='categorical_crossentropy',
↳metrics=['accuracy'])

# Print model summaries to compare parameter counts
print("Original Model Summary:")
original_model.summary()

```

Original Model Summary:

Model: "sequential\_4"

Layer (type)	Output Shape	
↳Param #		
conv2d_13 (Conv2D)	(None, 32, 32, 32)	
↳896		
max_pooling2d_13 (MaxPooling2D)	(None, 16, 16, 32)	
↳ 0		
conv2d_14 (Conv2D)	(None, 16, 16, 64)	
↳18,496		
max_pooling2d_14 (MaxPooling2D)	(None, 8, 8, 64)	
↳ 0		
conv2d_15 (Conv2D)	(None, 8, 8, 128)	
↳73,856		
max_pooling2d_15 (MaxPooling2D)	(None, 4, 4, 128)	
↳ 0		
flatten_4 (Flatten)	(None, 2048)	
↳ 0		
dense_12 (Dense)	(None, 128)	
↳262,272		
dense_13 (Dense)	(None, 64)	
↳8,256		
dense_14 (Dense)	(None, 10)	
↳650		

Total params: 364,426 (1.39 MB)

Trainable params: 364,426 (1.39 MB)

Non-trainable params: 0 (0.00 B)

```
[20]: print("\nModified Model Summary (with extra layer):")
      modified_model.summary()
```

Modified Model Summary (with extra layer):

Model: "sequential\_5"

Layer (type)	Output Shape	
↳Param #		
conv2d_16 (Conv2D)	(None, 32, 32, 32)	↳
↳896		
max_pooling2d_16 (MaxPooling2D)	(None, 16, 16, 32)	↳
↳ 0		
conv2d_17 (Conv2D)	(None, 16, 16, 64)	↳
↳18,496		
max_pooling2d_17 (MaxPooling2D)	(None, 8, 8, 64)	↳
↳ 0		
conv2d_18 (Conv2D)	(None, 8, 8, 128)	↳
↳73,856		
max_pooling2d_18 (MaxPooling2D)	(None, 4, 4, 128)	↳
↳ 0		
conv2d_19 (Conv2D)	(None, 4, 4, 256)	↳
↳295,168		
max_pooling2d_19 (MaxPooling2D)	(None, 2, 2, 256)	↳
↳ 0		
flatten_5 (Flatten)	(None, 1024)	↳
↳ 0		

```

dense_15 (Dense)                                     (None, 128)
↳131,200

dense_16 (Dense)                                     (None, 64)
↳8,256

dense_17 (Dense)                                     (None, 10)
↳650

```

Total params: 528,522 (2.02 MB)

Trainable params: 528,522 (2.02 MB)

Non-trainable params: 0 (0.00 B)

```

[21]: # Train both models
print("Training original model...")
history_original = original_model.fit(x_train, y_train, epochs=20,
↳batch_size=64, validation_data=(x_test, y_test))

print("\nTraining modified model with extra layer...")
history_modified = modified_model.fit(x_train, y_train, epochs=20,
↳batch_size=64, validation_data=(x_test, y_test))

# Final evaluation of both models
original_train_loss, original_train_accuracy = original_model.evaluate(x_train,
↳y_train, verbose=0)
original_test_loss, original_test_accuracy = original_model.evaluate(x_test,
↳y_test, verbose=0)

modified_train_loss, modified_train_accuracy = modified_model.evaluate(x_train,
↳y_train, verbose=0)
modified_test_loss, modified_test_accuracy = modified_model.evaluate(x_test,
↳y_test, verbose=0)

print(f"\nOriginal Model - Training Accuracy: {original_train_accuracy * 100:.
↳2f}%, Validation Accuracy: {original_test_accuracy * 100:.2f}%")
print(f"Modified Model - Training Accuracy: {modified_train_accuracy * 100:.
↳2f}%, Validation Accuracy: {modified_test_accuracy * 100:.2f}%")

```

Training original model...

Epoch 1/20

782/782 8s 7ms/step -

accuracy: 0.3675 - loss: 1.7192 - val\_accuracy: 0.5833 - val\_loss: 1.1669

Epoch 2/20  
782/782                    3s 4ms/step -  
accuracy: 0.6146 - loss: 1.0810 - val\_accuracy: 0.6660 - val\_loss: 0.9457  
Epoch 3/20  
782/782                    3s 4ms/step -  
accuracy: 0.6929 - loss: 0.8788 - val\_accuracy: 0.6971 - val\_loss: 0.8663  
Epoch 4/20  
782/782                    6s 5ms/step -  
accuracy: 0.7366 - loss: 0.7561 - val\_accuracy: 0.7114 - val\_loss: 0.8464  
Epoch 5/20  
782/782                    3s 4ms/step -  
accuracy: 0.7653 - loss: 0.6681 - val\_accuracy: 0.7060 - val\_loss: 0.8413  
Epoch 6/20  
782/782                    3s 4ms/step -  
accuracy: 0.7929 - loss: 0.5953 - val\_accuracy: 0.7247 - val\_loss: 0.8239  
Epoch 7/20  
782/782                    6s 4ms/step -  
accuracy: 0.8160 - loss: 0.5192 - val\_accuracy: 0.7338 - val\_loss: 0.7769  
Epoch 8/20  
782/782                    3s 4ms/step -  
accuracy: 0.8372 - loss: 0.4650 - val\_accuracy: 0.7081 - val\_loss: 0.8968  
Epoch 9/20  
782/782                    5s 4ms/step -  
accuracy: 0.8542 - loss: 0.4124 - val\_accuracy: 0.7441 - val\_loss: 0.8293  
Epoch 10/20  
782/782                    6s 5ms/step -  
accuracy: 0.8745 - loss: 0.3531 - val\_accuracy: 0.7388 - val\_loss: 0.8725  
Epoch 11/20  
782/782                    3s 4ms/step -  
accuracy: 0.8887 - loss: 0.3138 - val\_accuracy: 0.7350 - val\_loss: 0.8881  
Epoch 12/20  
782/782                    5s 4ms/step -  
accuracy: 0.9053 - loss: 0.2687 - val\_accuracy: 0.7264 - val\_loss: 1.0087  
Epoch 13/20  
782/782                    6s 5ms/step -  
accuracy: 0.9137 - loss: 0.2409 - val\_accuracy: 0.7304 - val\_loss: 1.0467  
Epoch 14/20  
782/782                    5s 4ms/step -  
accuracy: 0.9251 - loss: 0.2102 - val\_accuracy: 0.7312 - val\_loss: 1.0711  
Epoch 15/20  
782/782                    3s 4ms/step -  
accuracy: 0.9372 - loss: 0.1771 - val\_accuracy: 0.7285 - val\_loss: 1.1660  
Epoch 16/20  
782/782                    4s 5ms/step -  
accuracy: 0.9371 - loss: 0.1743 - val\_accuracy: 0.7397 - val\_loss: 1.1625  
Epoch 17/20  
782/782                    5s 4ms/step -  
accuracy: 0.9479 - loss: 0.1455 - val\_accuracy: 0.7301 - val\_loss: 1.2991

Epoch 18/20  
782/782 3s 4ms/step -  
accuracy: 0.9557 - loss: 0.1266 - val\_accuracy: 0.7257 - val\_loss: 1.3095  
Epoch 19/20  
782/782 3s 4ms/step -  
accuracy: 0.9558 - loss: 0.1253 - val\_accuracy: 0.7281 - val\_loss: 1.3956  
Epoch 20/20  
782/782 3s 4ms/step -  
accuracy: 0.9540 - loss: 0.1290 - val\_accuracy: 0.7231 - val\_loss: 1.4876

Training modified model with extra layer...

Epoch 1/20  
782/782 10s 9ms/step -  
accuracy: 0.3286 - loss: 1.8056 - val\_accuracy: 0.5919 - val\_loss: 1.1287  
Epoch 2/20  
782/782 6s 5ms/step -  
accuracy: 0.6081 - loss: 1.0840 - val\_accuracy: 0.6652 - val\_loss: 0.9514  
Epoch 3/20  
782/782 6s 5ms/step -  
accuracy: 0.7041 - loss: 0.8353 - val\_accuracy: 0.6772 - val\_loss: 0.9418  
Epoch 4/20  
782/782 4s 5ms/step -  
accuracy: 0.7581 - loss: 0.6899 - val\_accuracy: 0.7352 - val\_loss: 0.7802  
Epoch 5/20  
782/782 5s 5ms/step -  
accuracy: 0.7986 - loss: 0.5678 - val\_accuracy: 0.7387 - val\_loss: 0.7886  
Epoch 6/20  
782/782 5s 5ms/step -  
accuracy: 0.8321 - loss: 0.4727 - val\_accuracy: 0.7475 - val\_loss: 0.7569  
Epoch 7/20  
782/782 5s 5ms/step -  
accuracy: 0.8653 - loss: 0.3860 - val\_accuracy: 0.7434 - val\_loss: 0.8130  
Epoch 8/20  
782/782 6s 6ms/step -  
accuracy: 0.8911 - loss: 0.3119 - val\_accuracy: 0.7436 - val\_loss: 0.9003  
Epoch 9/20  
782/782 4s 5ms/step -  
accuracy: 0.9121 - loss: 0.2498 - val\_accuracy: 0.7461 - val\_loss: 0.8984  
Epoch 10/20  
782/782 5s 5ms/step -  
accuracy: 0.9313 - loss: 0.2011 - val\_accuracy: 0.7321 - val\_loss: 1.0650  
Epoch 11/20  
782/782 4s 5ms/step -  
accuracy: 0.9363 - loss: 0.1824 - val\_accuracy: 0.7425 - val\_loss: 1.0817  
Epoch 12/20  
782/782 4s 5ms/step -  
accuracy: 0.9512 - loss: 0.1411 - val\_accuracy: 0.7470 - val\_loss: 1.1716  
Epoch 13/20

```

782/782          4s 5ms/step -
accuracy: 0.9534 - loss: 0.1338 - val_accuracy: 0.7447 - val_loss: 1.2081
Epoch 14/20
782/782          4s 6ms/step -
accuracy: 0.9618 - loss: 0.1075 - val_accuracy: 0.7412 - val_loss: 1.2350
Epoch 15/20
782/782          4s 5ms/step -
accuracy: 0.9642 - loss: 0.1027 - val_accuracy: 0.7350 - val_loss: 1.2969
Epoch 16/20
782/782          4s 5ms/step -
accuracy: 0.9653 - loss: 0.1012 - val_accuracy: 0.7422 - val_loss: 1.3481
Epoch 17/20
782/782          6s 5ms/step -
accuracy: 0.9675 - loss: 0.0923 - val_accuracy: 0.7394 - val_loss: 1.2960
Epoch 18/20
782/782          5s 5ms/step -
accuracy: 0.9704 - loss: 0.0847 - val_accuracy: 0.7387 - val_loss: 1.5074
Epoch 19/20
782/782          4s 5ms/step -
accuracy: 0.9709 - loss: 0.0835 - val_accuracy: 0.7412 - val_loss: 1.5056
Epoch 20/20
782/782          6s 5ms/step -
accuracy: 0.9740 - loss: 0.0790 - val_accuracy: 0.7336 - val_loss: 1.3954

```

Original Model - Training Accuracy: 94.59%, Validation Accuracy: 72.31%  
Modified Model - Training Accuracy: 96.98%, Validation Accuracy: 73.36%

```

[22]: def plot_confusion_matrices(model1, model2, x_train, y_train, x_test, y_test):
        fig, axes = plt.subplots(2, 2, figsize=(20, 16))
        fig.suptitle('Confusion Matrices for Original and Modified Models',
        ↪fontsize=20)

        # Original Model - Training Data
        y_pred_train_orig = model1.predict(x_train)
        y_pred_classes_train_orig = np.argmax(y_pred_train_orig, axis=1)
        y_true_train = np.argmax(y_train, axis=1)
        cm_train_orig = confusion_matrix(y_true_train, y_pred_classes_train_orig)

        sns.heatmap(cm_train_orig, annot=True, fmt='d', cmap='Blues', ax=axes[0,
        ↪0], xticklabels=range(10), yticklabels=range(10))
        axes[0, 0].set_title('Original Model - Training Data')
        axes[0, 0].set_xlabel('Predicted')
        axes[0, 0].set_ylabel('Actual')

        # Original Model - Validation Data
        y_pred_test_orig = model1.predict(x_test)
        y_pred_classes_test_orig = np.argmax(y_pred_test_orig, axis=1)

```



```

y_true_test = np.argmax(y_test, axis=1)
cm_test_orig = confusion_matrix(y_true_test, y_pred_classes_test_orig)

sns.heatmap(cm_test_orig, annot=True, fmt='d', cmap='Blues', ax=axes[0, 1],
↳xticklabels=range(10), yticklabels=range(10))
axes[0, 1].set_title('Original Model - Validation Data')
axes[0, 1].set_xlabel('Predicted')
axes[0, 1].set_ylabel('Actual')

# Modified Model - Training Data
y_pred_train_mod = model2.predict(x_train)
y_pred_classes_train_mod = np.argmax(y_pred_train_mod, axis=1)
cm_train_mod = confusion_matrix(y_true_train, y_pred_classes_train_mod)

sns.heatmap(cm_train_mod, annot=True, fmt='d', cmap='Blues', ax=axes[1, 0],
↳xticklabels=range(10), yticklabels=range(10))
axes[1, 0].set_title('Modified Model - Training Data')
axes[1, 0].set_xlabel('Predicted')
axes[1, 0].set_ylabel('Actual')

# Modified Model - Validation Data
y_pred_test_mod = model2.predict(x_test)
y_pred_classes_test_mod = np.argmax(y_pred_test_mod, axis=1)
cm_test_mod = confusion_matrix(y_true_test, y_pred_classes_test_mod)

sns.heatmap(cm_test_mod, annot=True, fmt='d', cmap='Blues', ax=axes[1, 1],
↳xticklabels=range(10), yticklabels=range(10))
axes[1, 1].set_title('Modified Model - Validation Data')
axes[1, 1].set_xlabel('Predicted')
axes[1, 1].set_ylabel('Actual')

plt.tight_layout(rect=[0, 0, 1, 0.96])
plt.show()

```


```

[23]: plot_confusion_matrices(original_model, modified_model, x_train, y_train,
↳x_test, y_test)

```

1563/1563	3s 2ms/step
313/313	1s 2ms/step
1563/1563	3s 2ms/step
313/313	1s 2ms/step

Confusion Matrices for Original and Modified Models



```
[24]: def plot_training_history(history, title="Model"):
    # Plot training & validation accuracy values
    plt.figure(figsize=(14, 6))
    plt.suptitle(f'{title} - Training and Validation Metrics', fontsize=16)

    plt.subplot(1, 2, 1)
    plt.plot(history.history['accuracy'])
    plt.plot(history.history['val_accuracy'])
    plt.title('Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.xticks(np.arange(1, len(history.history['accuracy'])+1, 2))
    plt.legend(['Train', 'Validation'], loc='upper left')

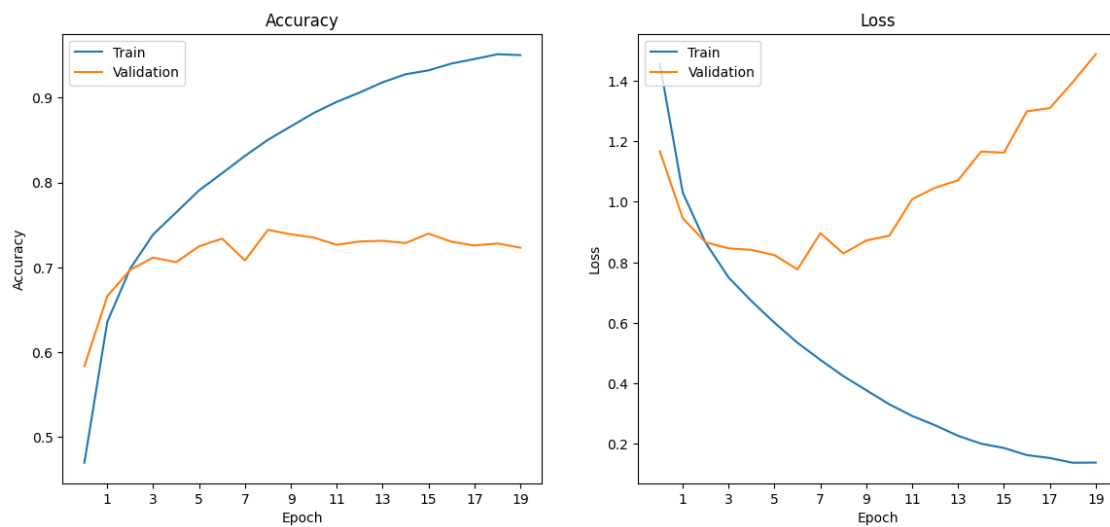
    # Plot training & validation loss values
    plt.subplot(1, 2, 2)
    plt.plot(history.history['loss'])
    plt.plot(history.history['val_loss'])
```

```
plt.title('Loss')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.xticks(np.arange(1, len(history.history['accuracy'])+1, 2))
plt.legend(['Train', 'Validation'], loc='upper left')


plt.show()
```

```
[25]: # Plot training history for both models
plot_training_history(history_original, title="Original Model")
plot_training_history(history_modified, title="Modified Model with Extra Layer")
```

Original Model - Training and Validation Metrics



Modified Model with Extra Layer - Training and Validation Metrics



## 2 Problem 2

```
[26]: # Load and preprocess CIFAR-10 data
(x_train, y_train), (x_test, y_test) = cifar10.load_data()

# Normalize the pixel values to be between 0 and 1
x_train, x_test = x_train.astype('float32') / 255.0, x_test.astype('float32') /
↪255.0
y_train = to_categorical(y_train, 10)
y_test = to_categorical(y_test, 10)

# Define the original CNN model without batch normalization
def create_original_cnn_model():
    model = models.Sequential()

    # First convolutional block
    model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same', ↪
↪input_shape=(32, 32, 3)))
    model.add(layers.MaxPooling2D((2, 2)))

    # Second convolutional block
    model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling2D((2, 2)))

    # Third convolutional block
    model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
    model.add(layers.MaxPooling2D((2, 2)))

    # Flatten and add dense layers
    model.add(layers.Flatten())
    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.Dense(64, activation='relu'))

    # Output layer
    model.add(layers.Dense(10, activation='softmax')) # 10 classes for CIFAR-10

    return model

# Define the original CNN model with batch normalization
def create_cnn_model_with_batch_norm():
    model = models.Sequential()

    # First convolutional block
```

```

    model.add(layers.Conv2D(32, (3, 3), activation='relu', padding='same',
↪input_shape=(32, 32, 3)))
    model.add(layers.BatchNormalization())
    model.add(layers.MaxPooling2D((2, 2)))

    # Second convolutional block
    model.add(layers.Conv2D(64, (3, 3), activation='relu', padding='same'))
    model.add(layers.BatchNormalization())
    model.add(layers.MaxPooling2D((2, 2)))

    # Third convolutional block
    model.add(layers.Conv2D(128, (3, 3), activation='relu', padding='same'))
    model.add(layers.BatchNormalization())
    model.add(layers.MaxPooling2D((2, 2)))

    # Flatten and add dense layers
    model.add(layers.Flatten())
    model.add(layers.Dense(128, activation='relu'))
    model.add(layers.BatchNormalization())
    model.add(layers.Dense(64, activation='relu'))
    model.add(layers.BatchNormalization())

    # Output layer
    model.add(layers.Dense(10, activation='softmax')) # 10 classes for CIFAR-10

    return model

```

```

[27]: # Instantiate both models
original_model = create_original_cnn_model()
original_model_bn = create_cnn_model_with_batch_norm()

# Compile both models
original_model.compile(optimizer='adam', loss='categorical_crossentropy',
↪metrics=['accuracy'])
original_model_bn.compile(optimizer='adam', loss='categorical_crossentropy',
↪metrics=['accuracy'])

```

```

/usr/local/lib/python3.10/dist-
packages/keras/src/layers/convolutional/base_conv.py:107: UserWarning: Do not
pass an `input_shape`/`input_dim` argument to a layer. When using Sequential
models, prefer using an `Input(shape)` object as the first layer in the model
instead.

```

```

    super().__init__(activity_regularizer=activity_regularizer, **kwargs)

```

```

[28]: # Train both models and record training time
import time

```

```

print("Training original model...")
start_time = time.time()
history_original = original_model.fit(x_train, y_train, epochs=20,
    ↪batch_size=64, validation_data=(x_test, y_test))
end_time = time.time()
original_training_time = end_time - start_time
print(f"Original model training completed in {original_training_time:.2f}
    ↪seconds.")

print("\nTraining original model with batch normalization...")
start_time = time.time()
history_original_bn = original_model_bn.fit(x_train, y_train, epochs=20,
    ↪batch_size=64, validation_data=(x_test, y_test))
end_time = time.time()
original_bn_training_time = end_time - start_time
print(f"Original model with batch normalization training completed in
    ↪{original_bn_training_time:.2f} seconds.")

```

Training original model...

Epoch 1/20

782/782 7s 6ms/step -

accuracy: 0.3347 - loss: 1.7947 - val\_accuracy: 0.5533 - val\_loss: 1.2140

Epoch 2/20

782/782 8s 4ms/step -

accuracy: 0.5932 - loss: 1.1288 - val\_accuracy: 0.6531 - val\_loss: 0.9881

Epoch 3/20

782/782 5s 4ms/step -

accuracy: 0.6836 - loss: 0.8978 - val\_accuracy: 0.6830 - val\_loss: 0.9038

Epoch 4/20

782/782 4s 5ms/step -

accuracy: 0.7292 - loss: 0.7706 - val\_accuracy: 0.7093 - val\_loss: 0.8360

Epoch 5/20

782/782 3s 4ms/step -

accuracy: 0.7592 - loss: 0.6844 - val\_accuracy: 0.7167 - val\_loss: 0.8283

Epoch 6/20

782/782 3s 4ms/step -

accuracy: 0.7927 - loss: 0.5953 - val\_accuracy: 0.7377 - val\_loss: 0.7625

Epoch 7/20

782/782 6s 4ms/step -

accuracy: 0.8137 - loss: 0.5381 - val\_accuracy: 0.7426 - val\_loss: 0.7707

Epoch 8/20

782/782 5s 4ms/step -

accuracy: 0.8365 - loss: 0.4649 - val\_accuracy: 0.7516 - val\_loss: 0.7567

Epoch 9/20

782/782 3s 4ms/step -

accuracy: 0.8594 - loss: 0.4066 - val\_accuracy: 0.7440 - val\_loss: 0.8107

Epoch 10/20

782/782                    3s 4ms/step -  
 accuracy: 0.8726 - loss: 0.3576 - val\_accuracy: 0.7460 - val\_loss: 0.8099  
 Epoch 11/20  
 782/782                    4s 5ms/step -  
 accuracy: 0.8888 - loss: 0.3185 - val\_accuracy: 0.7395 - val\_loss: 0.9044  
 Epoch 12/20  
 782/782                    5s 4ms/step -  
 accuracy: 0.9049 - loss: 0.2757 - val\_accuracy: 0.7388 - val\_loss: 0.9443  
 Epoch 13/20  
 782/782                    5s 4ms/step -  
 accuracy: 0.9150 - loss: 0.2426 - val\_accuracy: 0.7483 - val\_loss: 0.9468  
 Epoch 14/20  
 782/782                    4s 4ms/step -  
 accuracy: 0.9283 - loss: 0.2037 - val\_accuracy: 0.7400 - val\_loss: 0.9998  
 Epoch 15/20  
 782/782                    3s 4ms/step -  
 accuracy: 0.9362 - loss: 0.1821 - val\_accuracy: 0.7358 - val\_loss: 1.1077  
 Epoch 16/20  
 782/782                    3s 4ms/step -  
 accuracy: 0.9418 - loss: 0.1684 - val\_accuracy: 0.7370 - val\_loss: 1.2159  
 Epoch 17/20  
 782/782                    5s 6ms/step -  
 accuracy: 0.9500 - loss: 0.1451 - val\_accuracy: 0.7367 - val\_loss: 1.2360  
 Epoch 18/20  
 782/782                    8s 4ms/step -  
 accuracy: 0.9484 - loss: 0.1448 - val\_accuracy: 0.7335 - val\_loss: 1.3886  
 Epoch 19/20  
 782/782                    3s 4ms/step -  
 accuracy: 0.9530 - loss: 0.1311 - val\_accuracy: 0.7265 - val\_loss: 1.3616  
 Epoch 20/20  
 782/782                    5s 4ms/step -  
 accuracy: 0.9562 - loss: 0.1264 - val\_accuracy: 0.7320 - val\_loss: 1.3932  
 Original model training completed in 92.96 seconds.

Training original model with batch normalization...

Epoch 1/20  
 782/782                    13s 9ms/step -  
 accuracy: 0.4834 - loss: 1.4602 - val\_accuracy: 0.4436 - val\_loss: 1.7126  
 Epoch 2/20  
 782/782                    4s 5ms/step -  
 accuracy: 0.6974 - loss: 0.8734 - val\_accuracy: 0.6940 - val\_loss: 0.8784  
 Epoch 3/20  
 782/782                    6s 6ms/step -  
 accuracy: 0.7659 - loss: 0.6795 - val\_accuracy: 0.6957 - val\_loss: 0.8805  
 Epoch 4/20  
 782/782                    4s 5ms/step -  
 accuracy: 0.8128 - loss: 0.5520 - val\_accuracy: 0.7225 - val\_loss: 0.8184  
 Epoch 5/20

```

782/782          6s 5ms/step -
accuracy: 0.8470 - loss: 0.4398 - val_accuracy: 0.7223 - val_loss: 0.8515
Epoch 6/20
782/782          4s 5ms/step -
accuracy: 0.8808 - loss: 0.3489 - val_accuracy: 0.7451 - val_loss: 0.8040
Epoch 7/20
782/782          5s 5ms/step -
accuracy: 0.8993 - loss: 0.2917 - val_accuracy: 0.7352 - val_loss: 0.9167
Epoch 8/20
782/782          4s 5ms/step -
accuracy: 0.9260 - loss: 0.2203 - val_accuracy: 0.6903 - val_loss: 1.1059
Epoch 9/20
782/782          5s 5ms/step -
accuracy: 0.9378 - loss: 0.1810 - val_accuracy: 0.7003 - val_loss: 1.1159
Epoch 10/20
782/782          5s 5ms/step -
accuracy: 0.9443 - loss: 0.1606 - val_accuracy: 0.7254 - val_loss: 1.0847
Epoch 11/20
782/782          6s 5ms/step -
accuracy: 0.9523 - loss: 0.1402 - val_accuracy: 0.7341 - val_loss: 1.0977
Epoch 12/20
782/782          4s 5ms/step -
accuracy: 0.9565 - loss: 0.1224 - val_accuracy: 0.6926 - val_loss: 1.3914
Epoch 13/20
782/782          5s 5ms/step -
accuracy: 0.9581 - loss: 0.1191 - val_accuracy: 0.7357 - val_loss: 1.1727
Epoch 14/20
782/782          5s 5ms/step -
accuracy: 0.9652 - loss: 0.0982 - val_accuracy: 0.7436 - val_loss: 1.1696
Epoch 15/20
782/782          5s 5ms/step -
accuracy: 0.9716 - loss: 0.0845 - val_accuracy: 0.7512 - val_loss: 1.1373
Epoch 16/20
782/782          4s 6ms/step -
accuracy: 0.9694 - loss: 0.0876 - val_accuracy: 0.7448 - val_loss: 1.1655
Epoch 17/20
782/782          4s 5ms/step -
accuracy: 0.9707 - loss: 0.0846 - val_accuracy: 0.6923 - val_loss: 1.6226
Epoch 18/20
782/782          5s 5ms/step -
accuracy: 0.9712 - loss: 0.0830 - val_accuracy: 0.7434 - val_loss: 1.2910
Epoch 19/20
782/782          5s 5ms/step -
accuracy: 0.9752 - loss: 0.0694 - val_accuracy: 0.7393 - val_loss: 1.3204
Epoch 20/20
782/782          4s 5ms/step -
accuracy: 0.9756 - loss: 0.0712 - val_accuracy: 0.7419 - val_loss: 1.2885
Original model with batch normalization training completed in 105.26 seconds.

```



```
[29]: # Final evaluation of both models
original_train_loss, original_train_accuracy = original_model.evaluate(x_train,
    ↪y_train, verbose=0)
original_test_loss, original_test_accuracy = original_model.evaluate(x_test,
    ↪y_test, verbose=0)

original_bn_train_loss, original_bn_train_accuracy = original_model_bn.
    ↪evaluate(x_train, y_train, verbose=0)
original_bn_test_loss, original_bn_test_accuracy = original_model_bn.
    ↪evaluate(x_test, y_test, verbose=0)

# Print training and validation accuracy
print(f"\nOriginal Model - Training Accuracy: {original_train_accuracy:.4f},
    ↪Validation Accuracy: {original_test_accuracy:.4f}")
print(f"Original Model with Batch Normalization - Training Accuracy:
    ↪{original_bn_train_accuracy:.4f}, Validation Accuracy:
    ↪{original_bn_test_accuracy:.4f}")

# Print training times
print(f"Original Model Training Time: {original_training_time:.2f} seconds")
print(f"Original Model with Batch Normalization Training Time:
    ↪{original_bn_training_time:.2f} seconds")
```

Original Model - Training Accuracy: 0.9639, Validation Accuracy: 0.7320  
 Original Model with Batch Normalization - Training Accuracy: 0.9749, Validation  
 Accuracy: 0.7419  
 Original Model Training Time: 92.96 seconds  
 Original Model with Batch Normalization Training Time: 105.26 seconds

```
[30]: # Plot training and validation accuracy for both models on the same plot
import matplotlib.pyplot as plt

plt.figure(figsize=(10, 6))

# Plot training accuracy
plt.plot(history_original.history['accuracy'], label='Original Model - Training
    ↪Accuracy', linestyle='--')
plt.plot(history_original_bn.history['accuracy'], label='Original Model with BN
    ↪- Training Accuracy', linestyle='--')

# Plot validation accuracy
plt.plot(history_original.history['val_accuracy'], label='Original Model -
    ↪Validation Accuracy')
plt.plot(history_original_bn.history['val_accuracy'], label='Original Model
    ↪with BN - Validation Accuracy')
```

```
plt.title('Training and Validation Accuracy for Original Model and Model with_
↪Batch Normalization')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)
plt.show()
```

